

Assignment Set 1 — Defense Sheet

Computational Finance (WI4151), TU Delft — Prof. Fang Fang

Contents

1	The task (in my own words)	1
2	My approach and why	1
3	Key code sections (file + what it does + why I wrote it that way)	2
3.1	2b — 2bCF (2).py (sine density recovery, standard normal)	2
3.2	2d — 2dCF (2).py (sine-COS European put pricing)	2
3.3	2e — 2eCF (2).py (sine-COS CDF recovery)	3
3.4	3a — 3aCF (2).py (plain Monte Carlo quantile)	3
3.5	3b — 3bCF (2).py (importance sampling, mean shift)	4
3.6	3c — 3cCF (2).py (importance sampling, mean + variance shift)	4
4	Design decisions I must justify	4
5	Results and what they mean	5
6	The hardest question Fang could ask here — and my answer	5
7	“Show me in your code where you do X” — anticipated spots	6
8	Corrector feedback (from the graded PDF) — points lost and what to do	6

Scope: this sheet covers the **coded** sub-questions only (2b, 2d, 2e, 3a, 3b, 3c). The hand-derivation parts (1a, 1b, 1c, 2a, 2c) are in the PDF and are not repeated here, but Fang can jump from code to derivation, so the cross-references below matter.

1 The task (in my own words)

Assignment 1 has three blocks. The graded coded pieces sit in blocks 2 and 3.

- **Block 2 — the COS method, sine variant.** The paper builds COS from a Fourier-**cosine** expansion; the assignment asks me to redo it with a Fourier-**sine** expansion.
 - **2b:** program the sine version of eq. (11) and reproduce the density-recovery errors of Table 1 of the COS paper for the standard normal density.
 - **2d:** implement the sine-COS European **put** pricing formula (derived by hand in 2c) under Black-Scholes parameters $S_0 = 3$, $\sigma = 0.3$, $T = 1$, $K = 4$, $r = 0.03$, and report the absolute error vs the analytic BS put as a function of N .

- **2e**: extend the sine recovery one integration step further to recover the **CDF**, then verify it against the analytic log-normal CDF.
- **Block 3 — importance sampling for a tail quantile.** Portfolio of $N = 1000$ digital options; $L_T = \sum_i 1\{S_i(T) > K_i\}$ with a common Brownian factor W and idiosyncratic factors B_i . Estimate the 95% quantile of L_T .
 - **3a**: plain Monte Carlo.
 - **3b**: importance sampling, mean shift $W_T \sim N(1.5, T)$, reweight by p/q .
 - **3c**: importance sampling, mean + variance shift $W_T \sim N(1.5, 2T)$, reweight by p/q .

TODO (Bruno): In one sentence each, state the *point* of each sub-question — e.g. why 2b uses the standard normal density specifically, why 2e is “one extra integration step” beyond 2b/2c, and what the 3b-vs-3c comparison is meant to demonstrate about importance sampling.

2 My approach and why

What the code factually does (per block):

- **2b / 2e** build the sine series coefficients $F_k = (2/(b-a)) * \text{Im}\{ \phi(w_k) \exp(-i w_k a) \}$ with $w_k = k\pi/(b-a)$, then sum F_k against either $\sin(k \pi (x-a)/(b-a))$ (density, 2b) or the integrated sine $(1 - \cos(...))$ form (CDF, 2e).
- **2d** prices the put as $e^{-rT} * \sum_k \text{Im}\{ \phi(w_k) \exp(-i w_k a) \} * V_k$, where V_k is the analytic sine coefficient of the put payoff built from the ϕ_k/χ_k primitives.
- **3a/3b/3c** simulate $S_i(T)$ from the exact log-normal solution, count threshold crossings to form L_T , and take a (weighted, in 3b/3c) empirical 95% quantile.

TODO (Bruno): Justify, in your own words, the *modelling* choices that the code hard-codes: - 2b: why $[a, b] = [-10, 10]$ is a safe truncation for the standard normal. - 2d: why you used the **cumulant-based** range $[a, b] = c1 \pm L * \sqrt{c2 + \sqrt{c4}}$ with $L = 10$, and why that is the “correct” way to set the range rather than a fixed window. - 2e: why you switched to a **fixed** window $[a, b] = [-2, 2]$ for the CDF test instead of reusing the cumulant rule from 2d. - 3b/3c: why a mean shift of exactly 1.5, and why doubling the variance is the natural “second” scheme to test.

3 Key code sections (file + what it does + why I wrote it that way)

3.1 2b — 2bCF (2).py (sine density recovery, standard normal)

- 2bCF (2).py:4-5 set the truncation range $a = -10, b = 10$.
- 2bCF (2).py:6-7 set the evaluation grid $X_{test} = \text{arange}(-5, 6)$ (i.e. integers -5..5) and $N_{vals} = [4, 8, 16, 32, 64]$.
- 2bCF (2).py:10-12 $\text{normal_density}(x)$ returns the standard normal pdf $(1/\sqrt{2\pi}) \exp(-x^2/2)$.
- 2bCF (2).py:14-16 $\text{char_function}(w)$ returns $\exp(-0.5w^2)$, the characteristic function of $N(0, 1)$.
- 2bCF (2).py:18-33 $\text{sin_density}(x_vals, a, b, N)$:
 - :22 $k = \text{arange}(1, N + 1)$ — sine series starts at $k = 1$ (no $k = 0$ term).
 - :23 $w_k = k\pi/(b-a)$.
 - :26 $f_k = (2/(b-a)) * \text{Im}\{ \text{char}(w_k) * \exp(-i w_k a) \}$ — the sine coefficients F_k .
 - :30-31 for each x , forms $\sin(k \pi (x-a)/(b-a))$ and sums $f_k * \sin_terms$.

- 2bCF (2).py:39-43 loops over N , computes $\max|f_{\text{approx}} - f_{\text{true}}|$ and prints it.

TODO (Bruno): Explain *why* line 26 takes the **imaginary** part (tie it to the sine series and to Euler's formula in your 2a derivation), and why the coefficient prefactor is $2/(b-a)$ and not $1/(b-a)$.

3.2 2d — 2dCF (2).py (sine-COS European put pricing)

FACT / FORMATTING FLAG: this file is stored as a **single physical line** (the newlines were stripped). So I cannot give honest per-line numbers here — I reference it by **function name**. Fang may notice this; be ready to say it is a save/encoding artefact, not the logic.

- Parameters block: $S_0 = 3, \sigma = 0.3, T = 1, K = 4, r = 0.03; x_0 = \log(S_0/K); \mu = r - 0.5\sigma^2$.
- Cumulant range: $c_1 = x_0 + \mu T, c_2 = T\sigma^2, c_4 = 0, L = 10, a = c_1 - L\sqrt{c_2 + \sqrt{c_4}}, b = c_1 + L\sqrt{c_2 + \sqrt{c_4}}$ (the comment cites “eq 49”).
- $\phi_B(w, x, r, \sigma, T)$: characteristic function of $y = \ln(S_T/K)$, returns $\exp(i w \mu_y - 0.5 \text{var}_y w^2)$ with $\mu_y = x + (r - 0.5\sigma^2)T, \text{var}_y = \sigma^2 T$.
- $\phi_k(k, a, b, c, d)$: returns $(1/wk) * (\cos(wk(c-a)) - \cos(wk(d-a)))$ — the ψ_k primitive from 2c.
- $\chi_k(k, a, b, c, d)$: returns $\text{term}(d) - \text{term}(c)$ with $\text{term}(y) = e^y (\sin(wk(y-a)) - wk \cos(wk(y-a))) / (1 + wk^2)$ — the χ_k primitive from 2c.
- $V_k\text{put}(a, b, K, ks)$: returns $(2K/(b-a)) * (-\chi_k(ks, a, b, a, 0) + \phi_k(ks, a, b, a, 0))$, i.e. the payoff coefficient $V_k = (2K/(b-a))(\psi_k(a, 0) - \chi_k(a, 0))$.
- $\text{sin_put_price}(N, \dots)$: builds $\text{cf_part} = \text{Im}\{\phi_B(wk, x_0, \dots) \exp(-i wk a)\}$, multiplies by V_k , and returns $\exp(-rT) * \text{dot}(\text{cf_part}, V_k)$.
- $\text{bsput}(\dots)$: closed-form BS put for the reference price.
- The driver loops N in $[4, 8, 16, 32, 64, 128, 256, 512, 1024]$ and prints $|SIN - BS|$.

NAMING FLAG (be ready for this): in the report (2c) ψ_k is the integral of $\sin$ and χ_k the integral of $e^y \sin$. In this file the **function named** ϕ_k is the $\sin$ integral (the report's ψ_k) and the function named χ_k is the $e^y \sin$ integral. So ϕ_k here = ψ_k in the report. The numbers still match because $V_k\text{put}$ pairs them correctly, but the name ϕ_k collides with the characteristic function ϕ .

TODO (Bruno): Justify the *content* choices: why $c_4 = 0$ is acceptable here (what is the 4th cumulant of a Gaussian log-price?), why $x_0 = \log(S_0/K)$ is the right state variable, and why the discount factor is e^{-rT} rather than $e^{-r\Delta t}$ with some other Δt .

3.3 2e — 2eCF (2).py (sine-COS CDF recovery)

- 2eCF (2).py:5-8 parameters $r = 0.03, S_0 = 3, \sigma = 0.3, T = 1$.
- 2eCF (2).py:11-12 $\mu = (r - 0.5\sigma^2)T, \sigma_{\text{aux}} = \sigma \sqrt{T}$ — moments of $\ln(S_T/S_0)$.
- 2eCF (2).py:14-15 fixed range $a = -2, b = 2$.
- 2eCF (2).py:16-17 $X_{\text{test}} = \text{linspace}(-0.5, 0.5, 11), N_{\text{vals}} = [4, 8, 16, 32, 64, 128]$.
- 2eCF (2).py:19-23 $\text{char_function}(w) = \text{CF of } \ln(S_T/S_0): \exp(i w \mu - 0.5 \sigma_{\text{aux}}^2 w^2)$.
- 2eCF (2).py:25-29 $F_{\text{exact}}(z) = \text{norm.cdf}((z - \mu)/\sigma_{\text{aux}})$ — analytic normal CDF.
- 2eCF (2).py:31-46 $\text{sin_cdf}(z_{\text{vals}}, N, a, b)$:
 - :36-37 $ks = \text{arange}(1, N + 1), wk = ks \pi / (b - a)$.
 - :40 $im = \text{Imchar}(wk) \exp(-i w k a)$.

- :43 `cosine = cos( outer((z-a)/(b-a), ks) * pi )` — matrix of $\cos(k \pi (z-a)/(b-a))$.
- :46 returns $(2/\pi) * (\text{im}/ks) @ (1 - \text{cosine}).T$, i.e. $F(z) \approx (2/\pi) \sum_k (1/k) \text{Im}_k (1 - \cos(k \pi (z-a)/(b-a)))$.
- 2eCF (2).py:54-57 loop over N , print $\max |F_{\text{sin}} - F_{\text{true}}|$.

TODO (Bruno): Explain *why* the CDF formula carries a $1/k$ factor and a $2/\pi$ prefactor (where do they come from when you integrate the density term by term?), and why the $(1 - \cos(\dots))$ shape is the integrated form of $\sin(\dots)$.

3.4 3a — 3aCF (2).py (plain Monte Carlo quantile)

- 3aCF (2).py:5-12 parameters $N_{\text{stocks}}=1000$, $S_0=1$, $r=0.01$, $\text{sigma1}=0.8$, $\text{sigma2}=0.6$, $T=1$, $K=2$, $\text{alpha} = 1 - 0.95$.
- 3aCF (2).py:14 `drift = r - 0.5*(sigma1^2 + sigma2^2)`.
- 3aCF (2).py:16-31 `simulate_plainMC(paths)`:
 - :22 common factor $W_t \sim N(0, \text{sqrt}(T))$ (shape `paths`).
 - :23 idiosyncratic $B_i \sim N(0, \text{sqrt}(T))$ (shape `paths x N_stocks`).
 - :26 $S_T = S_0 \exp(\text{drift} * T + \text{sigma1} * W_t[:, \text{None}] + \text{sigma2} * B_i)$.
 - :29 $L_T = \text{sum}(S_T > K, \text{axis}=1)$ — count of in-the-money names per path.
- 3aCF (2).py:36 `np.random.seed(42)`.
- 3aCF (2).py:40-44 loops `paths` over $[1000..500000]$, $q = \text{np.quantile}(L_T, 1 - \text{alpha}) = 0.95$ quantile.
- 3aCF (2).py:47-55 semilog convergence plot.

TODO (Bruno): Justify why `np.std/np.random.normal(0, sqrt(T))` (a $N(0, T)$ draw) is the *exact* terminal law of W_T here, why no time-stepping is needed, and why the empirical quantile is taken at $1 - \text{alpha} = 0.95$ given the assignment's definition $P(L_T \geq q_{\text{alpha}}) \leq \text{alpha}$.

3.5 3b — 3bCF (2).py (importance sampling, mean shift)

- 3bCF (2).py:17-18 IS parameters $\mu_{\text{shift}} = 1.5$, $\mu_{\text{shiftsq}} = \mu_{\text{shift}}^2 / (2T)$.
- 3bCF (2).py:20-34 `weighted_quantile(values, weights, q)`:
 - :25-27 sort `values`, reorder `weights` to match.
 - :30 `cdf = cumsum(w_sorted)/sum(w_sorted)`.
 - :33-34 `idx = searchsorted(cdf, q)`, return `v_sorted[idx]`.
- 3bCF (2).py:43 $W_q \sim N(\mu_{\text{shift}}, \text{sqrt}(T))$ — common factor drawn under q .
- 3bCF (2).py:44 $B_i \sim N(0, \text{sqrt}(T))$ — idiosyncratic factors unchanged.
- 3bCF (2).py:47 $S_T = \exp(\text{drift} * T + \text{sigma1} * W_q[:, \text{None}] + \text{sigma2} * B_i)$ (note: no explicit S_0 , since $S_0 = 1$).
- 3bCF (2).py:50 $L_T = \text{sum}(S_T > K, \text{axis}=1)$.
- 3bCF (2).py:53 `ratio = exp(-(mu_shift/T)*W_q + mu_shiftsq)` — likelihood ratio p/q .
- 3bCF (2).py:56 `q_95 = weighted_quantile(L_T, ratio, 0.95)`.

TODO (Bruno): Derive the *ratio* on line 53 from p/q of two Gaussians and confirm it equals the report's $\exp(-1.5w + 1.125)$ at $T = 1$. Also justify *why only W is reweighted and not B_i* (the report says the B_i ratio is 1 — state precisely why).

3.6 3c — 3cCF (2).py (importance sampling, mean + variance shift)

- 3cCF (2).py:17-18 IS parameters $\mu_{\text{shift}} = 1.5$, $\text{var}_q = 2 * T$.
- 3cCF (2).py:20-34 `weighted_quantile` — identical structure to 3b.
- 3cCF (2).py:43 $W_q \sim N(\mu_{\text{shift}}, \text{sqrt}(\text{var}_q)) = N(1.5, 2T)$.

- `3cCF (2).py:44,47,50` same B_i , S_T , L_T construction as 3b.
- `3cCF (2).py:54` $\text{ratio} = \sqrt{2} \cdot \exp((-W_q^2 - 3 W_q + 2.25)/(4T))$ — the p/q for the mean+var shift.
- `3cCF (2).py:56` $q_{95} = \text{weighted_quantile}(L_T, \text{ratio}, 0.95)$.

TODO (Bruno): Derive the line-54 ratio from $p \sim N(0, T)$, $q \sim N(1.5, 2T)$ and confirm the `sqrt(2)` prefactor comes from the variance mismatch in the normalising constants. Then explain why this scheme is *worse* for the 95% quantile than 3b (the report’s argument about mass landing too far in the tail) — restate it as your own.

4 Design decisions I must justify

- **Sine vs cosine expansion.** The code implements the sine series throughout (coefficients via *Im...*). **TODO (Bruno):** why does the assignment want sine, and what changes structurally vs the cosine COS (the $k = 0$ half-weight term, the *Im* vs *Re*)?
 - **Truncation range, three different choices.** 2b uses $[-10, 10]$, 2d uses the cumulant rule with $L = 10$, 2e uses $[-2, 2]$. **TODO (Bruno):** justify each range and explain whether $[-2, 2]$ in 2e is wide enough for $N(\mu, \sigma_{aux}^2)$ with $\mu = -0.015$, $\sigma_{aux} = 0.3$ (it is $\sim \pm 6.7$ sigma — confirm and say so). This is exactly the kind of “is your range safe?” question Fang asks.
 - **$L = 10$ cumulant multiplier.** **TODO (Bruno):** the paper uses L around 8–10; justify your $L = 10$ and what error you’d see if L were too small (link to your Assignment 2 short-maturity / small- c_2 bug).
 - **np.quantile interpolation vs the *weighted_quantile* searchsorted rule.** 3a uses NumPy’s `np.quantile` (linear interpolation); 3b/3c use a custom step-function *searchsorted* quantile. **TODO (Bruno):** why two different quantile estimators, and does the mismatch bias the 3a-vs-3b/3c comparison?
 - **seed(42).** Used in 3a/3b/3c for reproducibility. **TODO (Bruno):** confirm this is only for reproducibility and does not, e.g., reuse the same W across schemes in a way that matters.
-

5 Results and what they mean

Reported in the PDF (factual, from the tables):

- **2b:** max density error falls $2.1e - 1(N = 4) \rightarrow 1.0e - 16(N = 64)$ — matches the paper’s Table 1 order of magnitude; spectral convergence.
- **2d:** $|SIN - BS|$ falls to $3.3e - 16$ by $N = 64$ and stays there; BS reference put = 0.9918191782.
- **2e:** max CDF error falls $1.37e - 1(N = 4) \rightarrow 3.67e - 11(N = 32)$ and then **plateaus** at $3.67e - 11$.
- **3a:** 95% quantile estimate settles near 590 as *paths* grows to $5e5$.
- **3b:** estimate converges to 587 and visibly stabilises with far fewer paths than 3a.
- **3c:** estimate ends near 589 but the path is noisier; the report argues it is *less* efficient than 3b for this quantile.

TODO (Bruno): - 2d/2e: explain the **plateau** (why does the error stop improving — round-off vs truncation-range error?). For 2e specifically, why does it plateau at $3.67e - 11$ and not at machine $1e - 16$ like 2d? This is the single most likely “your convergence stalled — why?” question. - 3a/3b/3c: are 590 / 587 / 589 *consistent* (do their confidence intervals overlap)? What is the “true” quantile and how confident are you?

6 The hardest question Fang could ask here — and my answer

Likely hardest: “Your 2e CDF error stalls at $3.67e-11$ while your 2d price reaches $1e-16$. Both use the sine-COS machinery. Why the difference, and is $3.67e-11$ a bug or expected?”

This bites because it sits exactly on the truncation-range / convergence theme Fang co-authored, and the two tables disagree by five orders of magnitude.

TODO (Bruno): Write the answer. Components you must cover: 1. The CDF formula integrates the density term by term, so its error inherits the **truncation-range** error $f(a), f(b) \neq 0$ at the endpoints — i.e. the residual comes from $[a, b] = [-2, 2]$ not being infinitely wide, not from too few terms N . (Confirm by saying what happens if you widen $[a, b]$.) 2. Why 2d reaches machine precision instead: the put payoff coefficients V_k decay and the BS density is extremely well captured by the cumulant range, so series-truncation error drops below round-off. 3. State plainly whether $3.67e-11$ is acceptable or whether widening $[-2, 2]$ would push it lower.

(Second candidate hardest: “Derive line 53 / line 54 likelihood ratios from scratch and tell me why the variance shift in 3c hurts.” — see the 3b/3c TODOs.)

7 “Show me in your code where you do X” — anticipated spots

- “...where you build the sine coefficients F_k .” -> 2bCF (2).py:26; same construction at 2eCF (2).py:40 and inside `sin_put_price(cfpart)` in 2dCF (2).py.
 - “...where you set the COS truncation range.” -> 2bCF (2).py:4-5 (fixed), cumulant block in 2dCF (2).py ($a = c1 - L * \sqrt{c2 + \sqrt{c4}}$), 2eCF (2).py:14-15 (fixed $[-2, 2]$).
 - “...where the characteristic function enters pricing.” -> ϕ_{BS} and $cfpart$ in 2dCF (2).py.
 - “...where you compute the payoff coefficients V_k .” -> `V_k_put` (with ϕ_k/χ_k) in 2dCF (2).py.
 - “...where you integrate the sine to get the CDF.” -> 2eCF (2).py:43-46 (the $(1 - \cosine)$ term).
 - “...where you get the exact terminal stock price.” -> 3aCF (2).py:26, 3bCF (2).py:47, 3cCF (2).py:47.
 - “...where you form L_T .” -> 3aCF (2).py:29, 3bCF (2).py:50, 3cCF (2).py:50.
 - “...where you apply the likelihood ratio.” -> 3bCF (2).py:53, 3cCF (2).py:54.
 - “...where you take the weighted quantile.” -> `weighted_quantile` in 3bCF (2).py:20-34 / 3cCF (2).py:20-34, called at :56.
 - “...where you take the plain quantile.” -> 3aCF (2).py:42 (`np.quantile(L_T, 1-alpha)`).
-

8 Corrector feedback (from the graded PDF) — points lost and what to do

Extracted from the grader's annotations in `Computational_Finance_Assignment_1(1)-3213313.pdf`. Two of these CONTRADICT earlier notes in this sheet — read carefully and rehearse the corrected version.

- **−0.2 (p5):** “Please also give the condition when $x \rightarrow -\infty$.” A missing limiting/boundary condition in the derivation. Add it.
- **−1.5 (2e, the $[a, b] = [-2, 2]$ choice) — DOCUMENTED COS-RANGE DEDUCTION:** “You manually chose $[a, b] = [-2, 2]$ in this question, which is not a very bad choice, but please choose the $[a, b]$ with the formula using cumulants as in the COS paper, which can give better error.” → **Reconcile with this sheet:** elsewhere I treated $[-2, 2]$ as “ $\sim \pm 6.7\sigma$, probably wide enough” and the $3.67e - 11$ plateau as merely “expected truncation-range error.” The grader’s verdict is sharper: $[-2, 2]$ was the wrong way to set the range and cost the **full 1.5 points**; the cumulant rule $[a, b] = c1 \pm L\sqrt{(c2 + \sqrt{c4})}$ (which you DID use in 2d) was expected, and the $3.67e - 11$ plateau is precisely the symptom of the hand-picked window. **This is the COS-range theme, graded — over-prepare it.** Correct answer: “I should have used the cumulant range in 2e as well; the fixed $[-2, 2]$ is what floored the CDF error at $3.67e - 11$.”
- **−1 (3b) and −1 (3c): IS-vs-MC conclusion not supported:** “From your numerical results of 3.a and 3.b, I cannot conclude that importance sampling is better than Monte Carlo... please test Monte Carlo and Importance Sampling several times and compare the variance/std of the quantile they get; then the variance/std of the Importance Sampling is evidently smaller.” → Action: don’t claim IS wins from a single run; run both **repeatedly** and compare the **variance/std of the estimated quantile**. Be ready to say this is how you’d demonstrate the variance reduction.
- **−0.2 (3c) — CONTRADICTS this sheet’s 3b-vs-3c claim:** “As long as p and q are pdfs for two normal distributions, then p/q is well defined. And if you compute the variance, you will see that in this example, the mean+var shift can usually get smaller variance than only the mean shift.” → **Reconcile:** this sheet (and the report) repeatedly says the **mean+var shift (3c) is less efficient** than the mean-only shift (3b). The grader says the opposite: the **mean+var shift usually gives smaller variance**. So do NOT rehearse “3c is worse.” The correct line: “3c (mean+var shift) typically achieves smaller variance than 3b; my report’s claim that it was worse was not well-supported — measured over repeated runs the mean+var shift is generally better.”

Action checklist: (1) add the $x \rightarrow -\infty$ condition; (2) say plainly the 2e range should have been the cumulant rule (this is the graded COS-range point); (3) demonstrate IS>MC via repeated-run variance comparison; (4) flip the 3b-vs-3c efficiency claim — mean+var shift is usually better.